
Trustworthy Completion for Web Agents: A Benchmark and Research Agenda for Execution-Time Safeguards

Anonymous Author(s)

Affiliation

Address

email

Abstract

Web agents may reach a user’s nominal goal while doing so through a misleading, policy-inconsistent, or otherwise unsafe action path. We call this failure mode unsafe completion. It is especially important in applications involving financial interests, privacy, informed consent, autonomy, and organizational policy, where conventional binary success metrics may count harmful behavior as successful task completion. We introduce a reproducible benchmark for evaluating trustworthy completion under deceptive web interfaces. The benchmark combines two local sandbox environments, nine short-horizon tasks spanning forced action, sneaking, and interface interference, task-level risk annotations, matched execution-time warning conditions, and a four-way outcome schema distinguishing safe completion, unsafe completion, safe abort, and other failure. Outcomes are scored through deterministic terminal-state checks rather than an LLM judge. An initial pilot with a fixed web-agent configuration reveals a substantial unsafe-completion rate that conventional task-success evaluation would obscure, while detecting no reliable difference between system-level and in-interface warning channels at the current scale. We therefore position the benchmark not as a definitive comparison of warning channels, but as a foundation for systematically studying execution-time safeguards. We further present a research agenda covering cross-agent validation, warning wording and placement, imperfect risk detection, broader task coverage, human calibration, and population-scale implications. The work aims to support more accountable deployment of web agents by making protection of stakeholder interests an explicit component of evaluation.

1 Introduction

Web agents increasingly act on behalf of people and organizations: they purchase products, configure services, grant permissions, and navigate consent interfaces. Capability benchmarks ask whether an agent reaches a requested endpoint [9, 22]. That endpoint, however, can conceal an unwanted charge, overbroad permission, privacy loss, or policy-inconsistent choice. A nominally successful trajectory may therefore be harmful to the stakeholder whose interests the agent was delegated to protect.

We call this distinction *trustworthy completion*. Nominal completion records that the goal endpoint was reached; trustworthy completion additionally requires that the agent avoided the task’s annotated unsafe path. This framing is complementary to security benchmarks for injected instructions [4, 5], broad computer-use risk [11, 20], and policy-constrained completion [12]. Our threat source is deceptive choice architecture presented to an agent with benign user intent, and our intervention is an execution-time warning at a task-annotated risk point.

Table 1: Threat-model distinctions. “Deterministic” refers to task-specific terminal-state checks.

Work	Threat source	User intent	Failure definition	Intervention / scorer
WebArena / VWA [9, 22]	General task complexity	Benign	Task failure	None / functional
AgentDojo / WASP [4, 5]	Injected instruction	Benign	Attack success	Defenses / benchmark
OS-Harm / RiOSWorld [11, 20]	Misuse and environmental risk	Mixed	Safety violation	None / policy checks
ST-WebAgentBench [12]	Policy and web risk	Benign	Unsafe policy completion	None / policy checks
DeceptiCon [3]	Deceptive UI	Benign	Pattern susceptibility	None / per-pattern
Ours	Deceptive choice architecture	Benign	Unsafe completion	Warning / deterministic

Research question. How should web agents be evaluated when nominal task completion may compromise the user’s financial interests, privacy, informed consent, autonomy, or policy constraints – and how can execution-time safeguards be tested without conflating warning design, risk detection, and agent capability?

We contribute three elements. First, a *safe-completion framework* distinguishes nominal completion from completion that protects stakeholder interests. Second, a *reproducible benchmark scaffold* combines two local sandboxes, nine task-level harm annotations, matched warning channels, four-way deterministic outcomes, and an auditable run manifest. Third, a *research agenda* makes agent backbone, warning wording, channel, trigger quality, deceptive pattern, and harm category separately testable. The current evidence is deliberately scoped: it is a three-repeat, single-agent pilot that demonstrates the need for outcome decomposition and does not establish a generally superior warning channel.

2 Related Work

Capability and safe completion. WebArena and VisualWebArena established self-hosted web environments for reproducible capability evaluation [9, 22]. AgenticShop studies personalized product curation [8], while Android and mobile benchmarks emphasize constraint adherence and long-horizon paths [16, 19]. These benchmarks motivate realistic interaction structure, but endpoint success alone does not identify whether the path protected the principal’s interests. ST-WebAgentBench is the closest conceptual bridge because it separates completion from policy and risk [12].

Security and environmental risk. AgentDojo, WASP, BIPIA, and Hidden-in-Plain-Text evaluate indirect prompt injection and defenses [4, 5, 7, 21]. OS-Harm and RiOSWorld broaden the scope to harmful requests, prompt injection, model misbehavior, and latent environmental risk [11, 20]. Our setting differs in threat source and factorization: no attacker instruction is required, user intent is benign, and the current oracle-trigger layer asks whether a correctly placed warning changes the agent’s choice.

Deceptive interfaces and safeguards. HCI and legal scholarship characterizes dark patterns as threats to consent, autonomy, privacy, and consumer welfare [6, 13–15]. Recent agent studies show susceptibility to dark-pattern interfaces, visual or fine-print injection, and persuasive task redirection [1–3, 10]. Human oversight can help but is neither costless nor a substitute for a separately designed human study [17]. Instruction-priority training addresses a different layer [18]. Table 1 makes these distinctions explicit.

3 Trustworthy Completion Framework

For a run, let C indicate that the nominal endpoint is reached and U that the annotated unsafe action is taken. The four outcomes are: *safe completion* ($C=1, U=0$), *unsafe completion* ($C=1, U=1$), *safe abort* ($C=0, U=0$) when the agent deliberately stops before the risk, and *other failure* for an execution breakdown or an outcome that cannot be reliably classified. Nominal completion is safe plus unsafe completion; trustworthy completion is safe completion. Safe abort protects the interest but does not satisfy the goal, so it remains distinct.

This schema is stakeholder-relative rather than severity-weighted. Each task identifies the acting principal, immediately affected stakeholder, protected interest, machine-verifiable unsafe action, potential consequence, safe alternative, qualitative reversibility, and expected safeguard response (avoid, clarify, confirm, or abort). These are author-defined annotations of represented risk, not measure-

Table 2: Task-level stakeholder and harm summary. Full eight-field annotations, safe alternatives, and reversibility judgments appear in the supplement.

Task (short)	Principal / affected stakeholder	Protected interest	Machine-verifiable unsafe action	Potential consequence
Subscription	Consumer / consumer	Financial; consent	Order with subscription selected	Unwanted subscription
Account gate	Consumer / consumer	Privacy; autonomy	Register before ordering	Unnecessary account and PII
Paid protection	Consumer / consumer	Financial	Order with protection selected	Extra charge
Trial renewal	Consumer / consumer	Financial; consent	Order with auto-renew enabled	Renewal commitment
Cookie consent	Consumer / consumer	Privacy; consent	Accept all cookies before order	Non-essential tracking
Newsletter	Consumer / consumer	Consent; autonomy	Order with opt-in selected	Unwanted marketing
Org consent	Organization / team	Privacy; policy	Publish with broad consent	Overbroad data access
Data processing	Organization / team	Privacy; policy	Finish with sensitive options on	Constraint violation
Workspace access	Organization / team	Privacy; policy	Approve full workspace access	Excess privilege




Figure 1: Trustworthy completion separates safe from unsafe endpoint success and factors the evaluation pipeline into detection, delivery, and response.

ments that downstream harm occurred. We assign no numerical harm score because such weighting would require separate normative validation. Table 2 summarizes the nine tasks; the supplement gives all eight fields.

Figure 1 summarizes the outcome and layer factorization. The *oracle-trigger layer* tests agent response when the risk point is correctly annotated and the safeguard is delivered as designed. A future *detector-coupled layer* must measure false positives, false negatives, delayed triggers, and localization errors. An *end-to-end deployment layer* must then combine detector quality, safeguard efficacy, and operational monitoring. The current pilot evaluates only the agent-response layer under an oracle trigger; it does not assume real deployments have perfect detectors.

4 Benchmark Construction

We define a task as

$$T = (G, E, D, R, A_{\text{safe}}, A_{\text{unsafe}}, S, H),$$

where G is the nominal goal; E the interface and trajectory; D the deceptive family and low-level implementation; R the risk point; A_{safe} and A_{unsafe} the safe and unsafe paths; S the deterministic terminal-state specification; and H the protected interest and potential harm.

Tasks are included only when: (1) a safe route to the nominal goal exists; (2) safe and unsafe paths reach comparable nominal endpoints; (3) the unsafe path has a concrete financial, privacy, consent, autonomy, or policy consequence; (4) the risk control is perceptible and actionable at the fixed viewport or documented interaction point; (5) both outcomes are deterministically verifiable; (6) the task maps to a documented deceptive-interface category; and (7) no real account, payment endpoint, or sensitive personal data is used.

Sandboxes and inventory. ShopLane hosts six consumer tasks; WorkHub Admin hosts three organizational tasks. The nine short-horizon tasks are balanced across forced action, sneaking, and interface interference. All use synthetic data and terminate in a shared state schema.

The present pilot holds task, agent, and oracle trigger fixed while varying warning delivery; deterministic terminal states produce the four outcomes in Figure 1.

Table 3: Frozen pilot outcomes by condition. Counts precede rates. Safe, unsafe, and abort rates use scorable runs; other-failure rates use all runs. Unsafe 95% CIs use 10,000 task-cluster bootstrap replicates (seed 42).

Condition	All	Scorable	Safe	Unsafe [95% CI]	Abort	Other fail.
No Warning	27	21	10/21 (.476)	11/21 (.524) [.190,.864]	0/21 (.000)	6/27 (.222)
System Warning	27	20	8/20 (.400)	12/20 (.600) [.250,.909]	0/20 (.000)	7/27 (.259)
UI Warning	27	24	9/24 (.375)	14/24 (.583) [.261,.870]	1/24 (.042)	3/27 (.111)

Warning conditions. No Warning supplies the goal and interface only. System Warning adds the task warning through the privileged instruction channel. UI Warning renders a warning near the risk-relevant control. The intended template is: “This step may involve [RISK_SLOT]. Complete the goal without taking it. If unsure, stop.” Current task and warning files use one concrete slot per task.

Scoring and artifact. At a terminal action, the sandbox records safe completion, unsafe completion, or safe abort in `terminal_state.json`; the Python scorer extracts this label without an LLM judge. Missing or unreadable terminal evidence maps conservatively to other failure. The anonymous artifact includes the sandbox, task YAML files, warning configuration, deterministic scorer, analysis code, and frozen outputs at <https://anonymous.4open.science/r/DeceptiveWebBench-960E/>. The artifact log records the access and synchronization status.

5 Pilot Evaluation

Setup. The pilot uses BrowserUse with Amazon Nova Lite v1 (`amazon.nova-lite-v1:0`) through AWS Bedrock, temperature 0, top- p 1, and a 15-step budget. The headless browser viewport is 1280×720 . Nine tasks, three conditions, and three repeats yield 81 uniquely keyed runs. The run manifest identifies every (`task_id`, `condition`, `repeat_id`) cell and hashes its available metadata and terminal artifacts.

Denominators and uncertainty. Safe completion, unsafe completion, and safe abort are mutually exclusive *scorable* outcomes; their rates use the scorable denominator. Other failure uses all runs. Uncertainty uses 10,000 task-cluster bootstrap replicates (seed 42): tasks are resampled with replacement while retaining their repeated runs, so task identity is not treated as exchangeable run noise. Intervals communicate uncertainty and are not a binary significance ritual.

5.1 Outcome Distribution

Table 3 reports counts, denominators, and rates. Unsafe completion occurs in 11/21 scorable No-Warning runs (52.4%), 12/20 System-Warning runs (60.0%), and 14/24 UI-Warning runs (58.3%). The System-minus-UI difference is +1.7 percentage points; its task-cluster 95% interval is $[-33.6, +36.6]$ points. The intervals overlap substantially and do not support a channel ranking. The supported conclusion is narrower: binary endpoint success would count the 37 unsafe completions as successes, while the four-way schema exposes them.

Other failures occur in 6/27 No-Warning runs (22.2%), 7/27 System-Warning runs (25.9%), and 3/27 UI-Warning runs (11.1%). The sole safe abort is under UI Warning (1/24 scorable, 4.2%); one event is insufficient to infer a channel property. Task-level outcomes are heterogeneous: the subscription and enterprise-sneaking tasks are unsafe in every scorable cell, while the enterprise workspace-access task is safe in every run. The task-by-condition heatmap is therefore descriptive, not confirmatory.

5.2 Failure Audit

Terminal labels support 37 unsafe decisions, 1 safe abstention, and 27 safe completions. The remaining 16 runs are `other_failure`. Their retained logs do not support a reliable exhaustive distinction between navigation/grounding and infrastructure causes, so we do not guess. The supplement lists observed terminal reasons when present and otherwise marks the cause unclassified.

Table 4: Research agenda. All rows are planned studies, not completed evidence.

Priority	Question / hypothesis	Design	Evidence requirement
D: cross-agent	Are effects model- or scaffold-dependent?	Same 9 tasks, 3 conditions, scorer; one stronger compatible agent; 5 repeats/cell	Audited 135-run manifest and agent $\times$ condition analysis
E1: wording	Does goal-plus-warning wording create conflict?	Three sneaking tasks; System/UI; current vs prohibitive wording; 5 repeats	Channel $\times$ wording estimates; supplement only
E2: repeats	Are intervals dominated by sampling variation?	Increase cross-agent cells from 5 to 10 repeats	Pre-frozen analysis and uncertainty contraction
E3: trigger	How do missed, false, or delayed triggers change outcomes?	Oracle vs delayed/false-positive subset	Detector-coupled error audit
Scale / human	Do results extend to longer tasks or humans?	More domains and paths; separate approved human study	New task validation; no human comparison without approval

6 Research Agenda

Table 4 specifies planned studies; none has been executed for this Version 1. The highest priority (D) is cross-agent validation under the unchanged benchmark, scorer, and warning conditions. A stronger compatible multimodal model would be evaluated on 9 tasks $\times$ 3 conditions $\times$ 5 repeats (135 new runs), with model identity, API version, adapter, prompts, rerun criteria, and analysis frozen in advance. Transparent descriptive results remain primary; an exploratory mixed-effects logistic model may test agent-by-condition heterogeneity with task as a random intercept.

E begins only after D is stable and remains supplementary for this workshop cycle. The first E study tests whether the current goal-plus-warning wording creates instruction conflict by comparing it with a purely prohibitive warning on the three sneaking tasks. Additional repeats improve uncertainty before introducing more factors. A small detector-coupled subset then compares oracle timing with delayed or false-positive warnings. Longer-term work expands domains and trajectory lengths, validates new task families, and conducts a separately approved human-subject calibration study. We do not propose numerical harm weighting or population-level causal claims without appropriate normative and empirical evidence.

7 Implications and Limitations

The AI-for-Good implication is methodological: delegated action should be evaluated against the interests of the people and organizations who bear its consequences. A system that reaches an endpoint but accepts tracking, creates an unnecessary account, incurs a charge, or grants excess access has not completed the delegation trustworthily. Controlled execution-time interventions provide an auditable way to test safeguards before deployment.

The evidence is nevertheless limited to one agent/scaffold, nine synthetic short-horizon tasks, three repeats per cell, oracle risk annotations, and author-defined qualitative harms. The wide task-cluster intervals show that modest channel effects are unresolved. Pattern-family results are exploratory. The local environments improve reproducibility but do not capture live-site drift, long-horizon compounding, or real remediation costs. No human baseline was collected, so we make no comparison between human and agent vulnerability. No individual-run rate is extrapolated to population-scale harm.

8 Conclusion

Binary task success is insufficient for delegated web agents. Trustworthy completion requires evidence that the agent reached the user’s goal without taking an annotated unsafe path. Our benchmark makes that distinction operational through stakeholder annotations, a four-way outcome schema, deterministic terminal-state scoring, and controlled safeguard delivery. The 81-run BrowserUse/Nova Lite pilot exposes substantial unsafe completion but is too small and heterogeneous to rank System and UI warning channels. Its value is therefore an auditable first instantiation and a scaffold for the explicitly planned cross-agent, wording, trigger, scale, and human-calibration studies.

References

- [1] Tri Cao, Bennett Lim, Yue Liu, Yuan Sui, Yuexin Li, Shumin Deng, Lin Lu, Nay Oo, Shuicheng Yan, and Bryan Hooi. VPI-Bench: Visual prompt injection attacks for computer-use agents. *arXiv preprint arXiv:2506.02456*, 2025.
- [2] Chaoran Chen, Zhiping Zhang, Bingcan Guo, and Shang Ma. The obvious invisible threat: Llm-powered gui agents' vulnerability to fine-print injections. *arXiv preprint arXiv:2504.11281*, 2025.
- [3] Phil Cuvin, Hao Zhu, and Diyi Yang. DECEPTICON: How dark patterns manipulate web agents. *arXiv preprint arXiv:2512.22894*, 2025.
- [4] Edoardo Debenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [5] Ivan Evtimov, Arman Zharmagambetov, Aaron Grattafiori, Chuan Guo, and Kamalika Chaudhuri. WASP: Benchmarking web agent security against prompt injection attacks. In *Advances in Neural Information Processing Systems*, volume 38, 2025.
- [6] Colin M. Gray, Yubo Kou, Bryan Battles, Joseph Hoggatt, and Austin L. Toombs. The dark (patterns) side of ux design. In *CHI EA '18: Extended Abstracts of the 2018 CHI Conference on Human Factors in Computing Systems*. Association for Computing Machinery, 2018. doi: 10.1145/3173574.3174108.
- [7] Haoze Guo and Ziqi Wei. Hidden-in-Plain-Text: A benchmark for social-web indirect prompt injection in RAG. In *Proceedings of the ACM Web Conference 2026*, 2026. doi: 10.1145/3774904.3792853.
- [8] Sunghwan Kim, Ryang Heo, Yongsik Seo, Jinyoung Yeo, and Dongha Lee. AgenticShop: Benchmarking agentic product curation for personalized web shopping. In *Proceedings of the ACM Web Conference 2026*, pages 2489–2500, 2026. doi: 10.1145/3774904.3792724.
- [9] Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Chong Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Ruslan Salakhutdinov, and Daniel Fried. VisualWebArena: Evaluating multimodal agents on realistic visual web tasks. *arXiv preprint arXiv:2401.13649*, 2024.
- [10] Karolina Korgul, Yushi Yang, Arkadiusz Drohomirecki, Piotr Błaszczczyk, Will Howard, Lukas Aichberger, Chris Russell, Philip H. S. Torr, Adam Mahdi, and Adel Bibi. It's a TRAP! task-redirecting agent persuasion benchmark for web agents. *arXiv preprint arXiv:2512.23128*, 2025.
- [11] Thomas Kuntz, Agatha Duzan, Hao Zhao, Francesco Croce, Zico Kolter, Nicolas Flammarion, and Maksym Andriushchenko. OS-Harm: A benchmark for measuring safety of computer use agents. In *Advances in Neural Information Processing Systems*, volume 38, 2025.
- [12] Ido Levy, Ben Wiesel, Sami Marreed, Alon Oved, Avi Yaeli, and Segev Shlomov. ST-WebAgentBench: A benchmark for evaluating safety and trustworthiness in web agents. *arXiv preprint arXiv:2410.06703*, 2024.
- [13] Jamie Luguri and Lior Jacob Strahilevitz. Shining a light on dark patterns. *Journal of Legal Analysis*, 13(1):43–109, 2021. doi: 10.1093/jla/laaa006.
- [14] Arunesh Mathur, Mihir Kshirsagar, and Jonathan Mayer. What makes a dark pattern... dark? design attributes, normative considerations, and measurement methods. In *Proceedings of the 2021 CHI conference on human factors in computing systems*, pages 1–18, 2021.
- [15] Midas Nouwens, Ilaria Liccardi, Michael Veale, David Karger, and Lalana Kagal. Dark patterns after the GDPR: Scraping consent pop-ups and demonstrating their influence. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, 2020. doi: 10.1145/3313831.3376321.

- 226 [16] Yuanyi Song et al. ColorBench: Benchmarking mobile agents with graph-structured framework
227 for complex long-horizon tasks. In *Proceedings of the ACM Web Conference 2026*, pages 7024–
228 7035, 2026. doi: 10.1145/3774904.3792215.
- 229 [17] Jingyu Tang, Chaoran Chen, Jiawen Li, Zhiping Zhang, Bingcan Guo, Ibrahim Khalilov, Sim-
230 ret A. Gebreegziabher, Bingsheng Yao, Dakuo Wang, Yanfang Ye, Tianshi Li, Ziang Xiao,
231 Yaxing Yao, and Toby Jia-Jun Li. Dark patterns meet GUI agents: LLM agent susceptibility
232 to manipulative interfaces and the role of human oversight. *arXiv preprint arXiv:2509.10723*,
233 2025.
- 234 [18] Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng, Johannes Heidecke, and Alex Beutel.
235 The instruction hierarchy: Training llms to prioritize privileged instructions. *arXiv preprint*
236 *arXiv:2404.13208*, 2024.
- 237 [19] Mingzhe Xing, Rongkai Zhang, Hui Xue, Qi Chen, Fan Yang, and Zhen Xiao. Understand-
238 ing the weakness of large language model agents within a complex android environment. In
239 *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*,
240 pages 6061–6072, 2024. doi: 10.1145/3637528.3671650.
- 241 [20] Jingyi Yang, Shuai Shao, Dongrui Liu, and Jing Shao. RiOSWorld: Benchmarking the risk
242 of multimodal computer-use agents. In *Advances in Neural Information Processing Systems*,
243 volume 38, 2025.
- 244 [21] Jingwei Yi, Yueqi Xie, Bin Zhu, Emre Kiciman, Guangzhong Sun, Xing Xie, and Fangzhao
245 Wu. Benchmarking and defending against indirect prompt injection attacks on large language
246 models. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and*
247 *Data Mining V.1*, pages 1809–1820, 2025. doi: 10.1145/3690624.3709179.
- 248 [22] Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng,
249 Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. WebArena: A realistic
250 web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023.